

# Support Vector Machine (SVM)

## Importing the libraries

```
In [10]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import datasets
%pylab inline
pylab.rcParams['figure.figsize'] = (10, 6)
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.

Populating the interactive namespace from numpy and matplotlib

## Importing the dataset

```
In [9]: from sklearn import datasets

new = datasets.load_diabetes()
print(new)
X = new.data[:, [2, 8]] # Feature 2: BMI, Feature 8: S5 (blood serum measure)
y = new.target
```

```

{'data': array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
                  0.01990749, -0.01764613],
                [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
                  -0.06833155, -0.09220405],
                [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
                  0.00286131, -0.02593034],
                ...,
                [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
                  -0.04688253,  0.01549073],
                [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
                  0.04452873, -0.02593034],
                [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
                  -0.00422151,  0.00306441]]), 'target': array([151., 75., 141., 20
6., 135., 97., 138., 63., 110., 310., 101.,
69., 179., 185., 118., 171., 166., 144., 97., 168., 68., 49.,
68., 245., 184., 202., 137., 85., 131., 283., 129., 59., 341.,
87., 65., 102., 265., 276., 252., 90., 100., 55., 61., 92.,
259., 53., 190., 142., 75., 142., 155., 225., 59., 104., 182.,
128., 52., 37., 170., 170., 61., 144., 52., 128., 71., 163.,
150., 97., 160., 178., 48., 270., 202., 111., 85., 42., 170.,
200., 252., 113., 143., 51., 52., 210., 65., 141., 55., 134.,
42., 111., 98., 164., 48., 96., 90., 162., 150., 279., 92.,
83., 128., 102., 302., 198., 95., 53., 134., 144., 232., 81.,
104., 59., 246., 297., 258., 229., 275., 281., 179., 200., 200.,
173., 180., 84., 121., 161., 99., 109., 115., 268., 274., 158.,
107., 83., 103., 272., 85., 280., 336., 281., 118., 317., 235.,
60., 174., 259., 178., 128., 96., 126., 288., 88., 292., 71.,
197., 186., 25., 84., 96., 195., 53., 217., 172., 131., 214.,
59., 70., 220., 268., 152., 47., 74., 295., 101., 151., 127.,
237., 225., 81., 151., 107., 64., 138., 185., 265., 101., 137.,
143., 141., 79., 292., 178., 91., 116., 86., 122., 72., 129.,
142., 90., 158., 39., 196., 222., 277., 99., 196., 202., 155.,
77., 191., 70., 73., 49., 65., 263., 248., 296., 214., 185.,
78., 93., 252., 150., 77., 208., 77., 108., 160., 53., 220.,
154., 259., 90., 246., 124., 67., 72., 257., 262., 275., 177.,
71., 47., 187., 125., 78., 51., 258., 215., 303., 243., 91.,
150., 310., 153., 346., 63., 89., 50., 39., 103., 308., 116.,
145., 74., 45., 115., 264., 87., 202., 127., 182., 241., 66.,
94., 283., 64., 102., 200., 265., 94., 230., 181., 156., 233.,
60., 219., 80., 68., 332., 248., 84., 200., 55., 85., 89.,
31., 129., 83., 275., 65., 198., 236., 253., 124., 44., 172.,
114., 142., 109., 180., 144., 163., 147., 97., 220., 190., 109.,
191., 122., 230., 242., 248., 249., 192., 131., 237., 78., 135.,
244., 199., 270., 164., 72., 96., 306., 91., 214., 95., 216.,
263., 178., 113., 200., 139., 139., 88., 148., 88., 243., 71.,
77., 109., 272., 60., 54., 221., 90., 311., 281., 182., 321.,
58., 262., 206., 233., 242., 123., 167., 63., 197., 71., 168.,
140., 217., 121., 235., 245., 40., 52., 104., 132., 88., 69.,
219., 72., 201., 110., 51., 277., 63., 118., 69., 273., 258.,
43., 198., 242., 232., 175., 93., 168., 275., 293., 281., 72.,
140., 189., 181., 209., 136., 261., 113., 131., 174., 257., 55.,
84., 42., 146., 212., 233., 91., 111., 152., 120., 67., 310.,
94., 183., 66., 173., 72., 49., 64., 48., 178., 104., 132.,
220., 57.] ), 'frame': None, 'DESCR': '.. _diabetes_dataset:\n\nDiabe
tes dataset\n-----\n\nTen baseline variables, age, sex, body mass
index, average blood\npressure, and six blood serum measurements were obtain

```

ed for each of  $n = 442$  diabetes patients, as well as the response of interest, a quantitative measure of disease progression one year after baseline.

**Data Set Characteristics:**

- Number of Instances: 442
- Number of Attributes: First 10 columns are numeric predictive values
- Target: Column 11 is a quantitative measure of disease progression one year after baseline

**Attribute Information:**

- age - age in years
- sex - sex
- bmi - body mass index
- bp - average blood pressure
- s1 - total serum cholesterol
- s2 - LDL, low-density lipoproteins
- s3 - HDL, high-density lipoproteins
- s4 - total cholesterol / HDL
- s5 - possibly log of serum triglycerides level
- s6 - glu, blood sugar level

Note: Each of these 10 feature variables have been mean centered and scaled by the standard deviation times the square root of  $n_{\text{samples}}$  (i.e. the sum of squares of each column totals 1).

Source URL: <https://www4.stat.ncsu.edu/~boos/var.select/diabetes.html>

For more information see: Bradley Efron, Trevor Hastie, Iain Johnstone and Robert Tibshirani (2004) "Least Angle Regression," Annals of Statistics (with discussion), 407-499. ([https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle\\_2002.pdf](https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf))

```
{
  'feature_names': ['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6'],
  'data_filename': 'diabetes_data_raw.csv.gz',
  'target_filename': 'diabetes_target.csv.gz',
  'data_module': 'sklearn.datasets.data'
}
```

## Exploratory Data Analysis

Place the diabetes data into a pandas dataframe

```
In [11]: new = datasets.load_diabetes()
X = new.data[:, [10, 20]] # Selecting two pixel features
y = new.target

# Create a DataFrame with column names representing pixel indices
digits_df = pd.DataFrame(X, columns=['pixel_10', 'pixel_20'])

# View the first 5 rows of the data
print(digits_df.head())

# Print the unique labels of the dataset
print('\nThe unique labels in this data are ' + str(np.unique(y)))
```

|   | pixel_10 | pixel_20 |
|---|----------|----------|
| 0 | 13.0     | 0.0      |
| 1 | 0.0      | 16.0     |
| 2 | 3.0      | 8.0      |
| 3 | 13.0     | 13.0     |
| 4 | 0.0      | 6.0      |

The unique labels in this data are [0 1 2 3 4 5 6 7 8 9]

## Splitting the dataset into the Training set and Test set

```
In [12]: from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3, random_state=42)

print('There are {} samples in the training set and {} samples in the test set'.format(X_train.shape[0], X_test.shape[0]))
```

There are 1257 samples in the training set and 540 samples in the test set

## Feature Scaling

```
In [13]: from sklearn.preprocessing import StandardScaler

sc = StandardScaler()

sc.fit(X_train)

X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)

print('After standardizing our features, the first 5 rows of our data now look like this:')
print(pd.DataFrame(X_train_std, columns=iris_df.columns).head())
```

After standardizing our features, the first 5 rows of our data now look like this:

|   | petal length (cm) | petal width (cm) |
|---|-------------------|------------------|
| 0 | 0.478495          | 0.141990         |
| 1 | 1.026780          | 1.428342         |
| 2 | -1.714644         | -0.983568        |
| 3 | -0.800836         | -0.661980        |
| 4 | 0.661257          | -1.144362        |

## Plot the original Data

```
In [16]: from matplotlib.colors import ListedColormap
import matplotlib.pyplot as plt
import numpy as np

# Assume X and y are defined
markers = ('s', 'x', 'o', '^', 'v', '>', '<', 'p', '*', 'D') # Add more if needed
colors = ('red', 'blue', 'lightgreen', 'orange', 'purple', 'cyan', 'brown', 'pink', 'gray')
cmap = ListedColormap(colors)

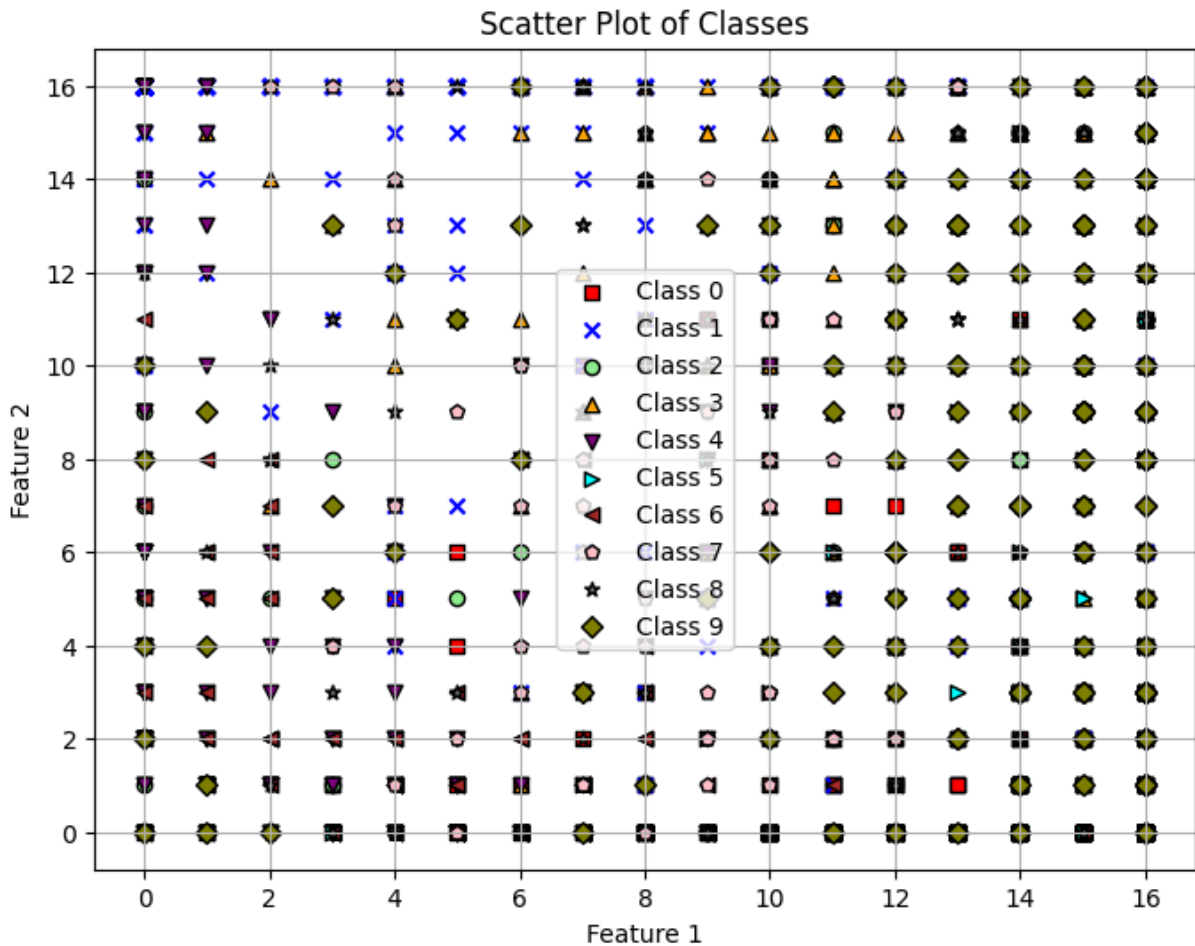
plt.figure(figsize=(8, 6))
for idx, cl in enumerate(np.unique(y)):
    plt.scatter(
        x=X[y == cl, 0],
        y=X[y == cl, 1],
        color=colors[idx % len(colors)], # Safe color cycling
        marker=markers[idx % len(markers)], # Safe marker cycling
        label=f'Class {cl}',
        edgecolor='k' if markers[idx % len(markers)] != 'x' else 'none' # avoid black outlines
    )

plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")
plt.title("Scatter Plot of Classes")
plt.legend()
plt.grid(True)
plt.show()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_34252\673965287.py:12: UserWarning: You passed a edgecolor/edgecolors ('none') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(
```



If we plot the original data, we can see that one of the classes is linearly separable, but the other two are not.

## Training the SVM model on the Training set

```
In [17]: from sklearn.svm import SVC
```

```
svm = SVC(kernel='rbf', random_state=0, gamma=.10, C=1.0)
svm.fit(X_train_std, y_train)
```

Out[17]:

```
SVC
SVC(gamma=0.1, random_state=0)
```

## Display the support Vectors of model

In [18]: `print("Support Vector for model are :",svm.support_vectors_)`

```
Support Vector for model are : [[ 0.84401829  0.46357804]
 [ 0.66125671 -0.66198024]
 [-0.61807438 -0.66198024]
 ...
 [ 1.02677988  0.302784  ]
 [ 0.47849512  0.78516612]
 [ 0.84401829  0.94596016]]
```

In [19]: `print("Number of support Vectors of each class 0 : - ",svm.n_support_[0])`  
`print("Number of support Vectors of each class 1 : - ",svm.n_support_[1])`  
`print("Number of support Vectors of each class 2 : - ",svm.n_support_[2])`

```
Number of support Vectors of each class 0 : - 133
Number of support Vectors of each class 1 : - 78
Number of support Vectors of each class 2 : - 120
```

In [20]: `print("Indices for support vectors are :",svm.support_)`

```
Indices for support vectors are : [ 14  18  39 ... 1212 1229 1232]
```

## Finding Accuracy of model on Test and Train Set

In [21]: `print('The accuracy of the svm classifier on training data is {:.2f} out of 1')`  
`print('The accuracy of the svm classifier on test data is {:.2f} out of 1')`

```
The accuracy of the svm classifier on training data is 0.39 out of 1
The accuracy of the svm classifier on test data is 0.35 out of 1
```

## Finding Accuracy of model on using confusion matrix

In [23]: `from sklearn.metrics import confusion_matrix`  
`cm = confusion_matrix(y_test, svm.predict(X_test_std))`  
`print(cm)`

```

[[ 6  0  0  3  0 21 10  1  0  4]
 [ 0 33  0  9  6  0  0  3  0  1]
 [ 0  2 30  6  3  3  3  2  0  4]
 [ 1  7 25 13  0  6  0  2  0  0]
 [ 0  7  0  3 19  2 15  2  0  0]
 [ 8  0  3  0  1 34  4  1  0  6]
 [ 7  0  0  0 13  0 40  0  0  0]
 [ 5  8  5 10  5  2  9  4  0  5]
 [ 0  6 14  8  3 13  4  2  0 11]
 [ 4  2 19  8  2  6  1  4  0 11]]

```

```

In [25]: from sklearn import metrics

# Predictions
y_pred = svm.predict(X_test_std)

# Core Metrics
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred, average='macro')
recall = metrics.recall_score(y_test, y_pred, average='macro') # aka sensit
f1 = metrics.f1_score(y_test, y_pred, average='macro')

# Specificity (manually computed per class from confusion matrix)
import numpy as np
cm = metrics.confusion_matrix(y_test, y_pred)
specificity_per_class = []
for i in range(len(cm)):
    tn = cm.sum() - (cm[i, :].sum() + cm[:, i].sum() - cm[i, i])
    fp = cm[:, i].sum() - cm[i, i]
    specificity = tn / (tn + fp)
    specificity_per_class.append(specificity)
specificity = np.mean(specificity_per_class) # macro-average specificity

# Print results
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall (Sensitivity): {recall:.4f}")
print(f"Specificity: {specificity:.4f}")
print(f"F1 Score: {f1:.4f}")

```

```

Accuracy: 0.3519
Precision: 0.2904
Recall (Sensitivity): 0.3502
Specificity: 0.9280
F1 Score: 0.3083

```

```

C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\skl
earn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is i
ll-defined and being set to 0.0 in labels with no predicted samples. Use `ze
ro_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

In [27]: print({
    "Accuracy": accuracy,
    "Precision": precision,
    "Sensitivity_recall": recall,
    "Specificity": specificity,

```

```
"F1_score": f1
}, end="")
```

```
{'Accuracy': 0.35185185185185186, 'Precision': 0.2904093805992233, 'Sensitivity_recall': 0.3502172576862249, 'Specificity': np.float64(0.927996260103083), 'F1_score': 0.308324979817366}
```

## Create the function for Visualizing Testing and Training model

```
In [28]: def versiontuple(v):
          return tuple(map(int, (v.split("."))))

def plot_decision_regions(X, y, classifier, test_idx=None, resolution=0.02):

    # setup marker generator and color map
    markers = ('s', 'x', 'o', '^', 'v')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.4, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())

    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0], y=X[y == cl, 1],
                    alpha=0.8, c=cmap(idx),
                    marker=markers[idx], label=cl)

plt.show()
```

## Visualising the Train set results

```
In [32]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

def plot_decision_regions(X, y, classifier, resolution=0.02):
    # Extended markers and colors to support more classes
    markers = ('s', 'x', 'o', '^', 'v', '>', '<', 'p', '*', 'D',
               'h', 'H', '+', '1', '2', '3', '4', '|', '_')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan', 'orange',
              'purple', 'pink', 'olive', 'brown', 'gold', 'teal',
              'coral', 'lime', 'navy', 'magenta', 'maroon', 'turquoise', 'ch
    cmap = ListedColormap(colors[:len(np.unique(y))])
```



```

# Mesh grid
x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                        np.arange(x2_min, x2_max, resolution))
Z = classifier.predict(np.c_[xx1.ravel(), xx2.ravel()])
Z = Z.reshape(xx1.shape)

# Plot contour
plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)
plt.xlim(xx1.min(), xx1.max())
plt.ylim(xx2.min(), xx2.max())

# Plot class data points
for idx, cl in enumerate(np.unique(y)):
    plt.scatter(x=X[y == cl, 0],
                y=X[y == cl, 1],
                color=colors[idx % len(colors)],
                marker=markers[idx % len(markers)],
                label=f'Class {cl}',
                edgecolor='k')

plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.title('Decision Regions')
plt.grid(True)
plt.show()


```

## Visualising the Test set results

```
In [33]: plot_decision_regions(X_test_std, y_test, svm)
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_34252\3420482309.py:29: UserWarning: You passed a edgecolor/edgecolors ('k') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == cl, 0],
```



In [ ]: