

Lab 10 : Model Selction

Import Commomn Libraries

```
In [79]: import itertools
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import NullFormatter
import pandas as pd
import numpy as np
import matplotlib.ticker as ticker
from sklearn import preprocessing
%matplotlib inline
```

About dataset

This dataset is about the chances of the Rain Today as per the given attributes like Temperature,Rainfall,Evaporation,Wind speed,Humidity and Pressure. It includes following fields:

| Field | Description |
|--------------|--|
| MinTemp | The Min Temperature of the day |
| MaxTemp | The Max Temperature of the day |
| Rainfall | The Rainfall Stats of that day |
| Evaporation | Evaporation Stats of that day |
| WindSpeed9am | Wind Speed stats till 9 am of that day |
| Humidity9am | Humidity stats till 9 am of that day |
| Pressure9am | Pressure stats till 9 am of that day |
| RainToday | The prediction of Raining at that day |

Import Training Dataset

```
In [80]: df = pd.read_csv(r"C:\Users\Admin\Downloads\weather_data.csv")
df.head()
```

| Out[80]: | MinTemp | MaxTemp | Rainfall | Evaporation | WindSpeed9am | Humidity9am |
|----------|---------|---------|----------|-------------|--------------|-------------|
| 0 | 9.4 | 30.6 | 3.4 | 8.0 | 25.9 | 63.4 |
| 1 | 23.8 | 29.4 | 4.8 | 7.1 | 6.9 | 86.6 |
| 2 | 18.3 | 15.8 | 0.9 | 4.9 | 34.9 | 30.0 |
| 3 | 15.0 | 25.2 | 2.9 | 2.1 | 24.5 | 30.4 |
| 4 | 3.9 | 26.4 | 2.5 | 7.0 | 6.3 | 64.5 |

EDA Steps

In [81]: `df.shape`

Out[81]: (1500, 8)

In [82]: `df.columns`

Out[82]: Index(['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'WindSpeed9am', 'Humidity9am', 'Pressure9am', 'RainToday'], dtype='object')

In [83]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MinTemp         1500 non-null   float64
1   MaxTemp         1500 non-null   float64
2   Rainfall        1500 non-null   float64
3   Evaporation     1500 non-null   float64
4   WindSpeed9am    1500 non-null   float64
5   Humidity9am     1500 non-null   float64
6   Pressure9am     1500 non-null   float64
7   RainToday       1500 non-null   object
dtypes: float64(7), object(1)
memory usage: 93.9+ KB
```

In [84]: `df.describe()`

| | MinTemp | MaxTemp | Rainfall | Evaporation | WindSpeed9am | H |
|--------------|-------------|-------------|-------------|-------------|--------------|---|
| count | 1500.000000 | 1500.000000 | 1500.000000 | 1500.000000 | 1500.000000 | |
| mean | 12.483533 | 30.012933 | 2.875733 | 6.995067 | 19.616400 | |
| std | 7.351179 | 8.680484 | 2.911895 | 2.898034 | 11.473366 | |
| min | 0.100000 | 15.000000 | 0.000000 | 2.000000 | 0.000000 | |
| 25% | 5.900000 | 22.700000 | 0.800000 | 4.500000 | 9.700000 | |
| 50% | 12.650000 | 30.150000 | 2.000000 | 6.900000 | 19.400000 | |
| 75% | 18.825000 | 37.625000 | 3.900000 | 9.500000 | 29.700000 | |
| max | 25.000000 | 45.000000 | 23.200000 | 12.000000 | 40.000000 | |

```
In [85]: df.describe(include = 'object')
```

```
Out[85]:
```

| | RainToday |
|--------|-----------|
| count | 1500 |
| unique | 2 |
| top | No |
| freq | 1062 |

Pre-processing Steps

Let's see how many of each class is in our data set

```
In [86]: df['RainToday'].value counts()
```

```
Out[86]: RainToday
No      1062
Yes      438
Name: count, dtype: int64
```

Pre-processing: Feature selection/extraction

Convert Categorical features to numerical values

Features before One Hot Encoding

```
In [87]: df[['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'WindSpeed9am', 'Humidity9am'
```

```
Out[87]:
```

| | MinTemp | MaxTemp | Rainfall | Evaporation | WindSpeed9am | Humidity9am |
|---|---------|---------|----------|-------------|--------------|-------------|
| 0 | 9.4 | 30.6 | 3.4 | 8.0 | 25.9 | 63.4 |
| 1 | 23.8 | 29.4 | 4.8 | 7.1 | 6.9 | 86.6 |
| 2 | 18.3 | 15.8 | 0.9 | 4.9 | 34.9 | 30.0 |
| 3 | 15.0 | 25.2 | 2.9 | 2.1 | 24.5 | 30.4 |
| 4 | 3.9 | 26.4 | 2.5 | 7.0 | 6.3 | 64.5 |

Use one hot encoding technique to cover categorical variables to binary variables and append them to the feature Data Frame

```
In [88]: Feature = df[['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'WindSpeed9am', 'Humidity9am', 'Pressure9am']]
Feature.head()
Feature.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MinTemp         1500 non-null   float64
1   MaxTemp         1500 non-null   float64
2   Rainfall        1500 non-null   float64
3   Evaporation     1500 non-null   float64
4   WindSpeed9am    1500 non-null   float64
5   Humidity9am     1500 non-null   float64
6   Pressure9am     1500 non-null   float64
dtypes: float64(7)
memory usage: 82.2 KB
```

Feature Selection

Let's define feature sets, X:

```
In [89]: X = Feature
X[:]
```

```
Out[89]:
```

| | MinTemp | MaxTemp | Rainfall | Evaporation | WindSpeed9am | Humidity9a |
|-------------|---------|---------|----------|-------------|--------------|------------|
| 0 | 9.4 | 30.6 | 3.4 | 8.0 | 25.9 | 63 |
| 1 | 23.8 | 29.4 | 4.8 | 7.1 | 6.9 | 86 |
| 2 | 18.3 | 15.8 | 0.9 | 4.9 | 34.9 | 30 |
| 3 | 15.0 | 25.2 | 2.9 | 2.1 | 24.5 | 30 |
| 4 | 3.9 | 26.4 | 2.5 | 7.0 | 6.3 | 64 |
| ... | ... | ... | ... | ... | ... | ... |
| 1495 | 16.7 | 41.0 | 2.8 | 8.8 | 3.1 | 61 |
| 1496 | 15.5 | 19.7 | 2.0 | 7.0 | 2.9 | 33 |
| 1497 | 11.6 | 24.3 | 0.4 | 9.7 | 34.0 | 98 |
| 1498 | 9.5 | 23.7 | 0.9 | 6.9 | 8.5 | 28 |
| 1499 | 21.6 | 41.1 | 1.6 | 3.5 | 26.8 | 94 |

1500 rows × 7 columns

What are our lables? Create Output Variable

```
In [90]: y = df['RainToday']
y[:]
d = {'No':0, 'Yes' : 1}
y = y.map(d)
```

Normalize Data

Data Standardization give data zero mean and unit variance (technically should be done after train test split)

```
In [91]: X= preprocessing.StandardScaler().fit(X).transform(X)
X[:]
```

```
Out[91]: array([[ -0.41960097,  0.06765319,  0.18010319, ...,  0.54785107,
                   0.20441611,  1.45735385],
 [ 1.53992184, -0.07063405,  0.66105015, ..., -1.10871051,
                   1.22432788,  0.09275376],
 [ 0.79149299, -1.6378895 , -0.67873068, ...,  1.33253813,
                   -1.26390515,  0.79555209],
 ...,
 [ -0.12022942, -0.65835484, -0.85049745, ...,  1.25406943,
                   1.73428377,  0.32116322],
 [ -0.40599317, -0.72749847, -0.67873068, ..., -0.96921059,
                   -1.32545155, -0.70375101],
 [ 1.2405503 ,  1.27766659, -0.4382572 , ...,  0.62631978,
                   1.5760216 , -1.62910214]])
```

Split the Data into Training and Testing Set

```
In [92]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, ra
print ('Train set:', X_train.shape, y_train.shape)
print ('Test set:', X_test.shape, y_test.shape)
```

Train set: (1200, 7) (1200,)

Test set: (300, 7) (300,)

Classification

Now, it is your turn, use the training set to build an accurate model. Then use the test set to report the accuracy of the model You should use the following algorithm:

- K Nearest Neighbor(KNN)
- Decision Tree
- Support Vector Machine
- Logistic Regression

__ Notice: __

- You can go above and change the pre-processing, feature selection, feature-extraction, and so on, to make a better model.
- You should use either scikit-learn, Scipy or Numpy libraries for developing the classification algorithms.
- You should include the code of the algorithm in the following cells.

K Nearest Neighbor(KNN)

Notice: You should find the best k to build the model with the best accuracy.\


```
In [93]: from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
Ks = 10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))

for n in range(1,Ks):
    knn1 = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat=knn1.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)

    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])

mean_acc
```

```
plt.plot(range(1,Ks),mean_acc,'g')
plt.fill_between(range(1,Ks),mean_acc - 1 * std_acc,mean_acc + 1 * std_acc,
plt.fill_between(range(1,Ks),mean_acc - 3 * std_acc,mean_acc + 3 * std_acc,
plt.legend(('Accuracy ', '+/- 1xstd','+/- 3xstd'))
plt.ylabel('Accuracy ')
plt.xlabel('Number of Neighbors (K)')
plt.tight_layout()
plt.show()
print( "The best accuracy was with", mean_acc.max(), "with k=", mean_acc.argmax())
```

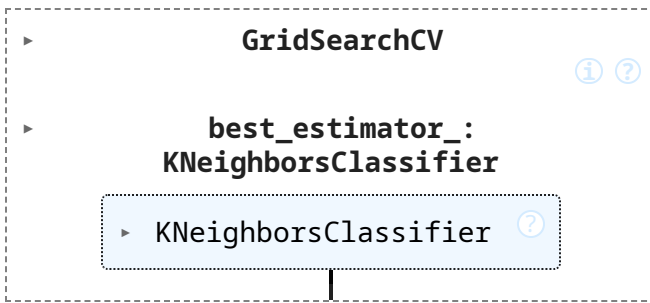


The best accuracy was with 0.67 with k= 8

Parameter Tuning using Grid Search Cv

```
In [94]: from sklearn.model_selection import GridSearchCV
k_range = list(range(1, 31))
weight_options = ['uniform', 'distance']
pow = [1,2]
param_grid = dict(n_neighbors=k_range,
                  weights=weight_options,
                  p = pow)
knn_gs = KNeighborsClassifier()
grid_k = GridSearchCV(knn_gs,
                    param_grid,
                    cv=10,
                    scoring='accuracy')
grid_k.fit(X_train, y_train)
```

```
Out[94]:
```



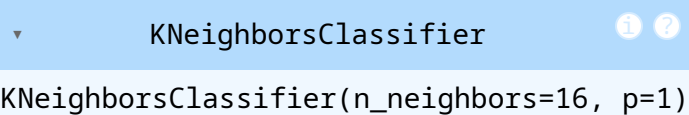
```
In [95]: print("Tuned Hyperparameters :", grid_k.best_params_)
print("Accuracy :",grid_k.best_score_)
```

Tuned Hyperparameters : {'n_neighbors': 30, 'p': 1, 'weights': 'uniform'}
Accuracy : 0.7133333333333334

```
In [96]: knn1 = KNeighborsClassifier(n_neighbors= 16, p = 1, weights = 'uniform')
```

```
In [97]: knn1.fit(X_train,y_train)
```

```
Out[97]:
```



```
In [98]: yhat = knn1.predict(X_test)
```

```
In [99]: from sklearn.metrics import jaccard_score
from sklearn.metrics import f1_score
from sklearn.metrics import log_loss
from sklearn.metrics import accuracy_score
a1 = jaccard_score(y_test,yhat,pos_label=1)
b1 = f1_score(y_test, yhat, average='weighted')
c1 = accuracy_score(y_test, yhat)
print('The jaccard_score of the KNN for k = 7 classifier on train data is {:.2f}')
print('The F1-score of the KNN for k = 7 classifier on train data is {:.2f}')
print('The Accuracy_score of the KNN for k = 7 classifier on train data is {:.2f}')
```

The jaccard_score of the KNN for k = 7 classifier on train data is 0.04
The F1-score of the KNN for k = 7 classifier on train data is 0.59
The Accuracy_score of the KNN for k = 7 classifier on train data is 0.69

Decision Tree

Parameter Tuning For Decision tree to find best tree

```
In [100]: from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeClassifier
param_grid = {'criterion': ['gini', 'entropy'],
               'max_features': ['auto', 'sqrt', 'log2'],
               'ccp_alpha': [0.1, .01, .001],
               'max_depth': [5, 6, 7, 8, 9],
```



```
    }  
    tree_clas = DecisionTreeClassifier(random_state=1)  
    grid_search = GridSearchCV(estimator=tree_clas, param_grid=param_grid, cv=5,  
    grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 90 candidates, totalling 450 fits

```
C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_validation.py:528: FitFailedWarning:
150 fits failed out of a total of 450.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

```
-----
----
150 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_validation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 1382, in wrapper
    estimator._validate_params()
  File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 436, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\utils\_param_validation.py", line 98, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of DecisionTreeClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_search.py:1108: UserWarning: One or more of the test scores are non-finite: [      nan  0.7125      0.7125      nan  0.7125
 0.7125
```

| | | | |
|----------------|------------|----------------|------------|
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.7125 | 0.7125 | nan 0.7125 | 0.7125 |
| nan 0.68833333 | 0.68833333 | nan 0.6925 | 0.6925 |
| nan 0.67916667 | 0.67916667 | nan 0.68333333 | 0.68333333 |
| nan 0.64166667 | 0.64166667 | nan 0.705 | 0.705 |
| nan 0.70166667 | 0.70166667 | nan 0.69166667 | 0.69166667 |
| nan 0.685 | 0.685 | nan 0.67583333 | 0.67583333 |

```
warnings.warn(
```

```
Out[100...] ▶ GridSearchCV ⓘ ?
  ▶ best_estimator_:
    DecisionTreeClassifier
      ▶ DecisionTreeClassifier ?
```

Find the best parameters

```
In [101...] print("Tuned Hyperparameters :", grid_search.best_params_)
print("Accuracy :", grid_search.best_score_)
```

```
Tuned Hyperparameters : {'ccp_alpha': 0.1, 'criterion': 'gini', 'max_depth':
5, 'max_features': 'sqrt'}
Accuracy : 0.7125
```

```
In [102...] Rain_Tree = grid_search.best_estimator_
print(Rain_Tree)
```

```
DecisionTreeClassifier(ccp_alpha=0.1, max_depth=5, max_features='sqrt',
random_state=1)
```

Train the best model using Training Data

```
In [103...] Rain_Tree.fit(X_train,y_train)
```

```
Out[103...] ▼ DecisionTreeClassifier ⓘ ?
DecisionTreeClassifier(ccp_alpha=0.1, max_depth=5, max_features='sq
rt',
random_state=1)
```

```
In [104...] predTree = Rain_Tree.predict(X_test)
```

```
In [105...] a2 = jaccard_score(y_test, predTree, pos_label=1)
b2 = f1_score(y_test, predTree, average='weighted')
c2 = accuracy_score(y_test, predTree)
print("The accuraccy of (Rain_tree) DecisionTrees's {:.2f} ".format(c2))
print('The jaccard_score of the (Rain_tree) DecisionTrees classifier on trai
print('The F1-score of the (Rain_tree) DecisionTrees classifier on train dat
```

```
The accuraccy of (Rain_tree) DecisionTrees's 0.69
The jaccard_score of the (Rain_tree) DecisionTrees classifier on train data
is 0.00
The F1-score of the (Rain_tree) DecisionTrees classifier on train data is 0.
56
```

Support Vector Machine

Parameter Tunning For SVM using GridSearchCV

```
In [106... from sklearn.svm import SVC
param_grid = {'C': [0.1, 1, 10, 100],
              'gamma': [1, 0.1, 0.01, 0.001],
              'kernel': ['rbf', 'sigmoid']}
grid_s = GridSearchCV(SVC(), param_grid, refit=True, verbose=2)
grid_s.fit(X_train, y_train)
```

Fitting 5 folds for each of 32 candidates, totalling 160 fits

```
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time=
0.1s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time=
0.1s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time=
0.1s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time=
0.1s
[CV] END .....C=0.1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=sigmoid; total time=
0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=sigmoid; total time=
```

0.0s
[CV] ENDC=0.1, gamma=0.01, kernel=sigmoid; total time=
0.0s
[CV] ENDC=0.1, gamma=0.01, kernel=sigmoid; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=rbf; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=rbf; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=rbf; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=rbf; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] ENDC=0.1, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=rbf; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=rbf; total time=
0.1s
[CV] ENDC=1, gamma=1, kernel=rbf; total time=
0.1s
[CV] ENDC=1, gamma=1, kernel=rbf; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=rbf; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] ENDC=1, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] ENDC=1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] ENDC=1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] ENDC=1, gamma=0.1, kernel=rbf; total time=
0.0s
[CV] ENDC=1, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] ENDC=1, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] ENDC=1, gamma=0.1, kernel=sigmoid; total time=

[illegible]

```
0.0s
[CV] END .....C=10, gamma=1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time=
0.1s
[CV] END .....C=10, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.1, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time=
0.2s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time=
0.2s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time=
0.3s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time=
0.2s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time=
0.2s
[CV] END .....C=10, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=10, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=10, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=10, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=10, gamma=0.01, kernel=sigmoid; total time=
0.0s
[CV] END .....C=10, gamma=0.001, kernel=rbf; total time=
0.0s
[CV] END .....C=10, gamma=0.001, kernel=rbf; total time=
```



[illegible]

```

0.0s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time=
0.5s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time=
0.6s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time=
0.7s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time=
0.4s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time=
0.5s
[CV] END .....C=100, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=100, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=100, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=100, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=100, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=100, gamma=0.01, kernel=sigmoid; total time=
0.1s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time=
0.1s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time=
0.0s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time=
0.1s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time=
0.1s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time=
0.1s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time=
0.1s
[CV] END .....C=100, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] END .....C=100, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] END .....C=100, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] END .....C=100, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] END .....C=100, gamma=0.001, kernel=sigmoid; total time=
0.0s
[CV] END .....C=100, gamma=0.001, kernel=sigmoid; total time=
0.0s

```

Out[106]...



In [107]...

```

print("Tuned Hyperparameters :", grid_s.best_params_)
print("Accuracy :", grid_s.best_score_)

```

```

Tuned Hyperparameters : {'C': 0.1, 'gamma': 1, 'kernel': 'rbf'}
Accuracy : 0.7125

```

```
In [108... svm = SVC(probability=True,C=0.1, gamma=1, kernel='sigmoid')
print(svm)
```

SVC(C=0.1, gamma=1, kernel='sigmoid', probability=True)

```
In [109... svm.fit(X_train, y_train)
```

```
Out[109... SVC
SVC(C=0.1, gamma=1, kernel='sigmoid', probability=True)
```

```
In [110... yhat_s = svm.predict(X_test)
```

```
In [111... a3 = jaccard_score(y_test, yhat_s,pos_label=1)
b3 = f1_score(y_test, yhat_s, average='weighted')
c3 = accuracy_score(y_test, yhat_s)
print('The accuracy of the svm classifier on training data is {:.2f} out of 1')
print('The accuracy of the svm classifier on test data is {:.2f} out of 1')
print('The jaccard_score of the SVM classifier on train data is {:.2f}'.format(a3))
print('The F1-score of the SVM classifier on train data is {:.2f}'.format(b3))
print('The accuracy-score of the SVM classifier on train data is {:.2f}'.format(c3))
```

The accuracy of the svm classifier on training data is 0.57 out of 1

The accuracy of the svm classifier on test data is 0.61 out of 1

The jaccard_score of the SVM classifier on train data is 0.09

The F1-score of the SVM classifier on train data is 0.57

The accuracy-score of the SVM classifier on train data is 0.61

Logistic Regression

Parameter Tunning using Grid Search CV

```
In [112... from sklearn.linear_model import LogisticRegression
parameters = {
    'penalty' : ['l1','l2'],
    'C'       : np.logspace(-3,3,7),
    'solver'  : ['newton-cg', 'lbfgs', 'liblinear'],
}
logreg = LogisticRegression()
clf = GridSearchCV(logreg,                                # model
                   param_grid = parameters,               # hyperparameters
                   scoring='accuracy',                    # metric for scoring
                   cv=10)                                  # number of folds
clf.fit(X_train,y_train)
```

```
C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_validation.py:528: FitFailedWarning:
140 fits failed out of a total of 420.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

```
-----
----
70 fits failed with the following error:
```

Traceback (most recent call last):

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_validation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 1389, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ~~~~~^
```

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear_model\_logistic.py", line 1193, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
    ~~~~~^
```

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear_model\_logistic.py", line 63, in _check_solver
    raise ValueError(
```

```
ValueError: Solver newton-cg supports only 'l2' or None penalties, got l1 penalty.
```

```
-----
----
70 fits failed with the following error:
```

Traceback (most recent call last):

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_validation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 1389, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ~~~~~^
```

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear_model\_logistic.py", line 1193, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
    ~~~~~^
```

```
File "C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear_model\_logistic.py", line 63, in _check_solver
    raise ValueError(
```

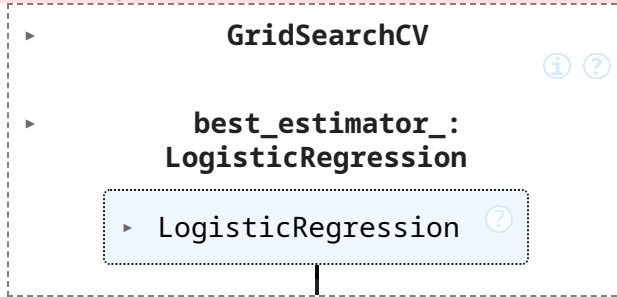
```
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1 penalty.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
```

```
C:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection\_search.py:1108: UserWarning: One or more of the test scores are non-finite: [ nan nan 0.7125 0.7125 0.7125 0.7125 nan nan 0.7125 0.7125
```

```
0.7125 0.7125 nan nan 0.7125 0.7125 0.7125 0.7125 nan nan
0.7125 0.7125 0.7125 0.7125 nan nan 0.7125 0.7125 0.7125 0.7125
nan nan 0.7125 0.7125 0.7125 0.7125 nan nan 0.7125 0.7125
0.7125 0.7125]
warnings.warn(
```

Out[112...



In [113...

```
print("Tuned Hyperparameters :", clf.best_params_)
print("Accuracy :",clf.best_score_)
```

```
Tuned Hyperparameters : {'C': np.float64(0.001), 'penalty': 'l1', 'solver':
'liblinear'}
Accuracy : 0.7125
```

In [114...

```
log_reg = clf.best_estimator_
log_reg.fit(X_train,y_train)
```

Out[114...

```

LogisticRegression ⓘ ?
LogisticRegression(C=np.float64(0.001), penalty='l1', solver='libli
near')
  
```

In [115...

```
yhat_l = log_reg.predict(X_test)
```

In [116...

```

a4 = jaccard_score(y_test, yhat_l,pos_label=1)
b4 = f1_score(y_test, yhat_l, average='weighted')
c4 = accuracy_score(y_test, yhat_l)

print('The jaccard_score of the logistic regression classifier on train data
print('The F1-score of the logistic regression classifier on train data is {
print('The accuracy_score of the logistic regression classifier on train dat
  
```

```

The jaccard_score of the logistic regression classifier on train data is 0.0
0
The F1-score of the logistic regression classifier on train data is 0.56
The accuracy_score of the logistic regression classifier on train data is 0.
69
  
```

Model Evaluation

In [117...

```

result_df=pd.DataFrame({'Model':['KNN','Decision Tree','SVM','Logistic Regre
'Jaccard Score' : [a1,a2,a3,a4],
'F1 Score' : [b1,b2,b3,b4],
'Accuracy Score':[c1,c2,c3,c4]})
  
```

In [118... `print(result_df)`

| | Model | Jaccard Score | F1 Score | Accuracy Score |
|---|---------------------|---------------|----------|----------------|
| 0 | KNN | 0.041237 | 0.585957 | 0.69 |
| 1 | Decision Tree | 0.000000 | 0.563432 | 0.69 |
| 2 | SVM | 0.093023 | 0.566884 | 0.61 |
| 3 | Logistic Regression | 0.000000 | 0.563432 | 0.69 |


using K-fold cross validation

```
In [119... from sklearn import model_selection
# prepare configuration for cross validation test harness
seed = 7
# prepare models
models = []
models.append(('KNN', knn1))
models.append(('DT', Rain_Tree ))
models.append(('SVM', svm))
models.append(('LR', log_reg))
# evaluate each model in turn
results = []
names = []
scoring = 'accuracy'
for name,model in models:
    kfold = model_selection.KFold(n_splits=10,
                                   random_state=seed,
                                   shuffle=True)
    cv_results = model_selection.cross_val_score(model,
                                                  X_train,
                                                  y_train,
                                                  cv=kfold,
                                                  scoring=scoring)

    results.append(cv_results)
    names.append(name)
for name,model in models:
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
# boxplot algorithm comparison
fig = plt.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

```
KNN: 0.712500 (0.041037)
DT: 0.712500 (0.041037)
SVM: 0.712500 (0.041037)
LR: 0.712500 (0.041037)
```

Algorithm Comparison



Plotting ROC_AUC Curce

```
In [120... #d = {'No':0,'Yes' : 1}  
#y_test = y_test.map(d)
```

```
In [121... from sklearn.metrics import roc_curve, roc_auc_score  
from sklearn.metrics import roc_curve, roc_auc_score  
  
# Instantiate the classifiers and make a list  
classifiers = [knn1,  
                Rain_Tree,  
                svm,  
                log_reg]  
model = ['KNN',  
         'Decision Tree',  
         'SVM',  
         'Logistic Regression', 'Naive Bayes']  
# Define a result table as a DataFrame  
result_table = pd.DataFrame(columns=['model', 'fpr', 'tpr', 'auc'])  
  
for cls in classifiers:  
    model = cls.fit(X_train, y_train)  
    yproba = model.predict_proba(X_test)[: , 1]  
  
    fpr, tpr, _ = roc_curve(y_test, yproba)  
    auc = roc_auc_score(y_test, yproba)
```

```

new_row = pd.DataFrame({
    'model': [cls.__class__.__name__],
    'fpr': [fpr],
    'tpr': [tpr],
    'auc': [auc]
})

result_table = pd.concat([result_table, new_row], ignore_index=True)

# Set name of the classifiers as index labels
result_table.set_index('model', inplace=True)

```

C:\Users\Admin\AppData\Local\Temp\ipykernel_76704\3286901965.py:30: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
result_table = pd.concat([result_table, new_row], ignore_index=True)
```

In [122... print(result_table)

```

                                                    fpr \
model
KNeighborsClassifier    [0.0, 0.0, 0.01932367149758454, 0.038647342995...
DecisionTreeClassifier                                [0.0, 1.0]
SVC                    [0.0, 0.0, 0.014492753623188406, 0.01449275362...
LogisticRegression                                [0.0, 1.0]

                                                    tpr \
model
KNeighborsClassifier    [0.0, 0.010752688172043012, 0.0430107526881720...
DecisionTreeClassifier                                [0.0, 1.0]
SVC                    [0.0, 0.010752688172043012, 0.0107526881720430...
LogisticRegression                                [0.0, 1.0]

                                                    auc
model
KNeighborsClassifier    0.485949
DecisionTreeClassifier    0.500000
SVC                      0.513168
LogisticRegression      0.500000

```

In [123... fig = plt.figure(figsize=(8,6))

```

for i in result_table.index:
    plt.plot(result_table.loc[i]['fpr'],
             result_table.loc[i]['tpr'],
             label="{}, AUC={:.3f}".format(i, result_table.loc[i]['auc']))

plt.plot([0,1], [0,1], color='orange', linestyle='--')

plt.xticks(np.arange(0.0, 1.1, step=0.1))
plt.xlabel("False Positive Rate", fontsize=15)

plt.yticks(np.arange(0.0, 1.1, step=0.1))
plt.ylabel("True Positive Rate", fontsize=15)

```



```
plt.title('ROC Curve Analysis', fontweight='bold', fontsize=15)
plt.legend(prop={'size':13}, loc='lower right')

plt.show()
```

